

COMENIUS UNIVERSITY, BRATISLAVA
FACULTY OF MATHEMATICS, PHYSICS AND INFORMATICS

BASE MANIFOLD MESHES FROM SKELETONS

MASTER THESIS

2013

Bc. Michal Piovarči

COMENIUS UNIVERSITY, BRATISLAVA
FACULTY OF MATHEMATICS, PHYSICS AND INFORMATICS

BASE MANIFOLD MESHES FROM SKELETONS

MASTER THESIS

Study programme: Computer Science
Study field: 9.2.1 Computer Science, informatics
Department: Department of Computer Science
Supervisor: RNDr. Martin Madaras

Bratislava, 2013

Bc. Michal Piovarči



Comenius University in Bratislava
Faculty of Mathematics, Physics and Informatics

THESIS ASSIGNMENT

Name and Surname:

Study programme:

Field of Study:

Type of Thesis:

Language of Thesis:

Title:

Aim:

Supervisor:

Department:

Assigned:

Approved:

Guarantor of Study Programme

Student

Supervisor



Comenius University in Bratislava
Faculty of Mathematics, Physics and Informatics

THESIS ASSIGNMENT

Name and Surname:

Study programme:

Field of Study:

Type of Thesis:

Language of Thesis:

Title:

Aim:

Supervisor:

Department:

Assigned:

Approved:

Guarantor of Study Programme

.....
Student

.....
Supervisor

Acknowledgement

I would like to thank my supervisor RNDr. Martin Madaras for his help and advices.

Abstract

We have implemented a method, that generates base manifold mesh from an input skeleton, based on Skeleton to Quad Dominant Mesh (SQM) algorithm, which converts skeletons to meshes composed mainly from quadrilaterals. Each node in skeleton has assigned a sphere. SQM algorithm first creates branch node polyhedrons for each sphere corresponding to a branch node. These polyhedrons are bridged with quadrilaterals to create the final base mesh. We have extended the algorithm to support generation of meshes from cyclic skeletons. We have also generalized skeleton nodes to ellipsoids instead of spheres. Finally, we extended the algorithm to generate meshes from linear skeletons without branching and from skeletons, which root node is not a branch node. The generated base mesh is tessellated on GPU for better visual results.

KEYWORDS: skeleton, convert, base mesh, manifold

Abstrakt

Implementovali sme algoritmus, ktorý generuje základné manifoldné meshe zo vstupnej kostry, založený na algoritme Skeleton to Quad Dominant Mesh (SQM), ktorý konvertuje kostry na meshe zložené prevažne zo štvoruholníkov. Každý vrchol kostry má priradenú guľu. Algoritmus SQM najprv pre každý vrchol v ktorom sa kostra vetví vygeneruje polyhedrón, aproximujúci guľu príslušného vetviaceho sa vrcholu. Neskôr sú tieto polyhedróny spojené štvoruholníkmi, čím sa vytvorí základný mesh. Rozšírili sme SQM algoritmus tak, aby generoval základné meshe aj z cyklických kostier. Tiež sme zovšeobecnilí vstupnú kostru tak, aby bolo možné zadávať elipsoidy namiesto gulí pre každý vrchol. Nakoniec sme rozšírili algoritmus tak aby vedel generovať základné meshe aj z lineárnych kostier a z kostier, ktorých koreň nie je vetviaci sa uzol. Vygenerovaný základný mesh teselujeme na grafickej karte, aby sme dosiahli lepšie vizuálne výsledky.

Kľúčové slová: kostra, konverzia, manifold, základný mesh

Preamble

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nunc tristique, sem et feugiat ornare, lorem eros mattis odio, et tempus lectus ipsum nec ante. Phasellus interdum nunc ut sapien semper porttitor. Nam mi erat, faucibus in fermentum eu, varius eu velit. Integer egestas iaculis varius. In pulvinar, ligula eget adipiscing suscipit, nisl ipsum aliquet arcu, eget tristique felis leo vitae magna. Nulla et magna sed justo accumsan ultrices a in leo. Suspendisse tincidunt malesuada leo, eget rhoncus ipsum fringilla at. Integer et tortor vitae nisl fermentum vestibulum. Fusce eu dui neque, a egestas nunc. Vivamus condimentum mi non arcu lacinia et aliquam risus euismod. Nunc ut risus nec elit luctus aliquet et sit amet magna. Vestibulum vehicula enim eget erat fermentum a lacinia purus varius.

Duis tempus sem sit amet elit accumsan ultricies. Curabitur a nibh ante, vitae pharetra nulla. Suspendisse non risus elit, in aliquam felis. Maecenas suscipit placerat commodo. Vivamus et molestie odio. Quisque ut augue mi. Quisque aliquam luctus est, ac dignissim ante adipiscing eget. Quisque volutpat, sem vitae placerat condimentum, nunc lorem malesuada leo, sit amet pretium nisi felis nec lorem. Pellentesque nisi ipsum, vestibulum sed lacinia sed, condimentum a turpis.

In posuere convallis lectus vel hendrerit. Cras suscipit mi risus. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Donec ante nunc, cursus ac vulputate at, bibendum eget nisi. Nunc eget nunc sed massa blandit posuere id vel quam. Duis bibendum orci vel ligula tempor condimentum. Nulla pharetra tortor at risus dignissim fringilla. Nullam ac massa et nibh auctor vestibulum quis vitae ligula. Suspendisse ultrices eros sit amet lectus dictum dapibus. Sed congue, turpis nec aliquam fermentum, diam nisi cursus nibh, id vulputate massa tellus sit amet turpis.

Contents

Acknowledgement	v
Abstract	vi
Abstrakt	vii
Preamble	viii
Introduction	1
1 Related Work	2
1.1 B-Mesh algorithm	3
1.2 Skeleton to Quad Dominant Mesh algorithm	5
1.3 Skeleton-based and interactive 3D modelling	7
2 Proposed solution	9
2.1 Skeleton straightening	9
2.1.1 Skinning matrices	10
2.2 BNP generation	11
2.3 BNP refinement	13
2.3.1 Smoothing	15
2.4 BNP joining	17
2.5 Final vertex placement	19
2.6 Ellipsoid Nodes	20
2.7 Tessellation	20
2.8 Capsule Ending	21
2.9 Linear Skeletons	22

<i>CONTENTS</i>	x
2.10 Root That Is Not a Branch Node	23
2.11 Cyclic Skeletons	24
3 Implementation	27
3.1 Programming Language, IDE, Tools, and Libraries	27
3.2 Model	29
3.3 View	29
3.4 Controller	29
3.5 GPU Shaders	29
A T_EX	30

List of Figures

1.1	B-Mesh sweeping and stitching illustration	3
1.2	Steps of SQM algorithm	6
1.3	Spherical nodes radii	8
2.1	Skeleton straightening	10
2.2	Rotation estimation from reference pose	11
2.3	BNP generation process	12
2.4	Obtuse triangle problem	13
2.5	BNP refinement	14
2.6	LIE smoothing schemes	16
2.7	BNP joining process	18
2.8	BNP joining pairing	19
2.9	Ellipsoid nodes	20
2.10	Tessellation and smoothing	21
2.11	Capsule generation	22
2.12	Linear base mesh	23
2.13	Delaunay triangulation closing cyclic mesh	25
2.14	Cyclic skeleton generation overview	26

List of Tables

Introduction

Insert introduction.

Chapter 1

Related Work

Generating base meshes from skeletons aids in modelling and rigging of more complex meshes. The most notable algorithms are B-Mesh[6] by Ji et al. and Skeleton to Quad Dominant Mesh[9] (SQM) by J. A. Bærentzen et al. These algorithms are capable of generating quad-dominant manifold meshes with good edge flow. Generated meshes are convenient because thanks to quad dominance and good edge flow they are easily skinned. Since the mesh is generated from skeleton we also implicitly know how much each bone affects vertices of the mesh during animation. We will also discuss some results from Michal McDonnell's master thesis Skeleton-based and interactive 3D modelling[2]. In particular the generation of capsule endings in SQM algorithm as we also wanted our algorithm to be capable of generating capsules.

B-Mesh algorithm and SQM algorithm present two different approaches to generation of base meshes from skeletons. B-Mesh algorithm firstly generates mesh for the limbs of the skeletons. These limb meshes are later joined together. On the other SQM algorithm firstly creates polyhedrons for branch nodes of the input skeleton. These polyhedrons are later joined with tubes consisting solely from quadrilaterals.

In this chapter we will describe both algorithms with their advantages and disadvantages. This serves to show why we have based our implementation on SQM algorithm.

1.1 B-Mesh algorithm

Input B-Mesh algorithm takes as input a skeleton with a set of spheres or ellipsoids. Each node of the input skeleton has assigned a sphere or ellipsoid that represents its local geometry. Auxiliary spheres can be added to more precisely affect the resulting geometry of generated base mesh.

The algorithm consist of five steps:

Step 1: Sphere generation - new spheres are generated along the bones of the skeleton.

Step 2: Sweeping - generation of mesh for skeletons limbs.

Step 3: Stitching - joining of limb meshes at branch nodes.

Step 4: B-Mesh evolution - subdivision and smoothing of generated mesh.

Step 5: B-Mesh fairing - fairing to improve edges flow.

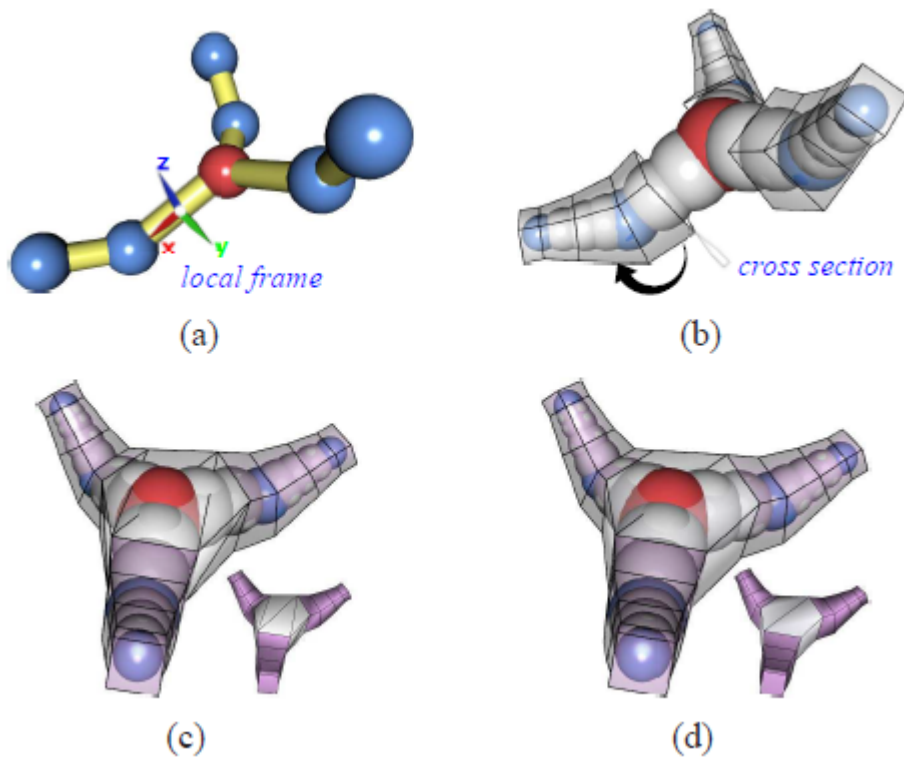


Figure 1.1: B-Mesh sweeping and stitching illustration. (a) local coordinates of a bone; (b) sweeping step; (c) stitching step; (d) after stitching simplification; Source [6]

Sphere generation Additional spheres are generated along the bones of the input skeleton. The distance between two spheres is defined by sampling step and the radius of each sphere is interpolated from radius's corresponding to the bones nodes. These generated spheres are used to refine the generated mesh and to calculate scalar field need in B-Mesh fairing step.

Sweeping For each bone its local coordinate system is generated as can be seen in Figure 1.1 (a). Limb mesh generation starts at branch nodes. For each new limb a rectangle aligned with limbs corresponding bones local axis is generated. Its points are then translated along the bones forming said limb and rotated around the connection nodes. After each translation the new points are connected with previous set of points in order to form a tube consisting of quadrilaterals. The resulting tubes can be seen in Figure 1.1 (b).

Stitching Limb meshes generated in sweeping step are now stitched together. This is done by generating a convex hull from all the points corresponding to each branch nodes. These points are the beginning of each limb mesh tube. The result can be seen in Figure 1.1 (c). Generated convex hull is composed from triangles. To simplify them into quadrilaterals a score between each pair of triangles is computed. The score measures how close to a plane is each pair of triangles. The results of stitching simplification are shown in Figure 1.1 (d).

B-Mesh evolution Catmull-Clark subdivision is used to smooth the stitched mesh. However after smoothing the mesh shrinks and deviates from the spheres. A scalar field is generated and used to evolve the mesh. Each vertex of the stitched mesh is translated according to the scalar field and its evolution speed. This means the further away is the vertex from its corresponding sphere the more it is attracted to it. In this phase the auxiliary spheres are used to manipulate the scalar field and thus the final shape of the mesh. This step can be repeated multiple times to further smooth the mesh.

B-Mesh fairing After the evolution step certain edges may not be aligned with their principal directions. New vertex positions are calculated by iterative minimization of a function.

Conclusion The biggest problem with B-Mesh algorithm for our use is its iterative nature. The number of iterations is an input parameter and we wanted to avoid input parameters, so that the base mesh generation is as automatic as possible. Without the evolution step B-Mesh

approximates input geometry worse than SQM. This is caused by the stitching phase which creates rectangular geometry at branch nodes, instead of the input spherical or elliptical geometry. Because of this evolution phase can not be avoided if we want to approximate the input as much as possible. B-Mesh algorithm is also running slower and produces more triangles than SQM according to [9]. The auxiliary spheres and ellipsoid nodes are interesting additions that were not present in SQM. But we wanted to modify the generated base mesh geometry directly on GPU during rendering so they are not advantageous to our intended use. B-mesh algorithm can generate capsules but it needs several evolution steps to precisely match capsules geometry. The algorithm looks like it can be naturally extended to support cyclic skeletons. Sweeping steps occurs only on limb paths and stitching occurs directly at branch nodes. Evolution and fairing moves vertices directly so it is independent from cycles in the input skeleton.

1.2 Skeleton to Quad Dominant Mesh algorithm

Input The input to SQM algorithm is a skeleton. Similarly to B-mesh for each of skeletons nodes a sphere is defined to represent local geometry. Contrary to B-mesh algorithm SQM does not support ellipsoid nodes. SQM algorithm has also an extra limitation that the root of the skeleton has to be a branch node.

The algorithm consist of four steps and one preprocessing step:

Preprocessing: Skeleton straightening - serves to simplify step number 3 of the algorithm.

Step 1: BNP generation - generation of branch node polyhedrons (BNPs).

Step 2: BNP refinement - subdivisions of BNPs.

Step 3: Creating the tubular structure - bridging of BNPs.

Step 4: Vertex placement - reverting straightened mesh to its original pose.

Skeleton straightening This is a preprocessing step of the algorithm that simplifies the generation of tubular structures. For each connection node its child is rotated, so that the edge between connection node and its child is parallel with the edge between connection node and its parent. This is useful because during step 3 the algorithm needs to only generate straight tubes and does not need to take rotation into account.




Figure 1.2: Steps of SQM algorithm. (a) the input skeleton; (b) generated BNPs; (c) refined BNPs; (d) BNPs bridges with quadrilateral tubes; Source [9]

BNP generation A branch node polyhedron(BNP) is a polyhedron assigned to a branch node. Vertices of a BNP correspond to a set of points that are generated by intersecting the sphere assigned to a branch node with each edge connected to said branch node. We will call this vertices intersection vertices. To form a BNP intersection vertices are triangulated. After that each triangle is split into six triangles by inserting one vertex in the middle of each triangle and in the middle of each of the edges of the triangle. These vertices are then projected back onto the sphere associated with a branch node. This projection is needed because if the intersection vertices are coplanar the generated polyhedron would have no volume. The result of this step can be seen in Figure 1.2 (b).

BNP refinement During step 3 of the algorithm BNPs connected via path are bridged with tubes consisting solely from quadrilaterals. This is done by connecting the one-rings of two corresponding intersection vertices with faces. To ensure that we can use only quadrilaterals the one-rings need to have the same valency. So each BNP is refined so that the valency of two intersection nodes lying on the same path are equal. This can be seen in Figure 1.2 (c).

Creating the tubular structure After previous step of the algorithm connected BNPs can be joined by a tube formed by quadrilaterals. The tube is divided into segments. Each of the segments corresponds to a connection node. Vertices corresponding to a certain connection node are projected onto its corresponding sphere. Leaf nodes are terminated with a triangle which central vertex corresponds to the leaf nodes vertex. The result is illustrated in Figure

1.2 (d).

Vertex placement The base mesh is now finished. All that remains is to reverse the rotations used to straighten the input skeleton. After final vertex placement the resulting mesh is smoothed with three iterations of Laplacian smoothing and attraction scheme.

Conclusion Each of SQMs steps are one pass algorithms except the final smoothing and attraction scheme application. But since SQM generates branch node geometry by generating a polyhedron it better resembles the geometry of the input skeleton even without smoothing. Cyclic meshes can be problematic to implement, because in the refinement step of the algorithm a cyclic mesh may cause an infinite loop of refining. SQM produces smaller number of triangles because in the joining phase all triangles corresponding to branch nodes are removed and their former vertices are used to generate connecting tubes. On the other hand during the stitching phase B-Mesh introduces triangles into the mesh. SQM algorithm was also capable of generating cycles from two symmetric leaf nodes, that were close to each other. However these cycles can not be created explicitly. Another limitation is that cycles have to lie on input skeletons axis of symmetry.

1.3 Skeleton-based and interactive 3D modelling

Michal Mc Donells master thesis studies the capabilities of SQM algorithm. It aims to extend SQM algorithm through pre-processing of SQMs input skeleton and post-processing of the output mesh. Altering SQM algorithm itself was out of scope of the thesis. Among the improvements to the algorithm were cone, cube, sphere and hemisphere leaf nodes, cube connection and branch nodes as well as concavities and branch node root requirement. The most important parts for us are spherical leaf nodes and branch node root requirement.

Spherical leaf node Original SQM algorithm was not able to create spheres at leaf nodes. This also limited its ability to generate capsules that we wanted to implement in our algorithm. In the thesis a solution was presented that a spherical node would be represented by several connection nodes. The first connection node would have the same radius as the desired sphere and each subsequent node would have its radius decreased until the desired number of subdivisions would be reached. We can see the process in Figure 1.3. An arc




Figure 1.3: Left: arc subdivided at fixed intervals; Right: radii of newly inserted nodes sampled from an arc on the left; Source [2]

with equal radius as the original spherical node would be sampled at fixed intervals. With enough samples the resulting mesh would approximate a sphere. However in the original SQM algorithm after Laplacian smoothing the smoothed mesh did not approximate the input sphere.

Conclusion The approach to create spherical nodes presented in the thesis seems suitable for our needs. Since we would not implement Laplacian smoothing we would not have problems with deformation of the generated sphere. On the contrary this approach is beneficial to our needs, since we wanted to smooth the generated mesh in tessellation shaders and keeping spherical meshes generated without post-processing can potentially ease the work-flow.

Root branch node requirement One of the limitations of SQM algorithm is that the root node has to be a branch node. The proposed solution to this limitation was to generate pseudo paths from the root to ensure that it always will have at least three child nodes. For example, if the root has only two childes the new node would be generated in a direction perpendicular to a line connecting the intersection of roots associated sphere and edges from root to its childes.

Conclusion This pseudo path generation would disturb the edge flow as it would introduce new unnecessary edges. We believe that a better approach would be to re-root the tree and replace its former root with a node with at least three childes. Linear skeletons without branching and skeletons consisting only from a single node would become special cases that would be handled separately. This would allow us to avoid the requirement without altering the expected edge flow of the input skeleton.

Chapter 2

Proposed solution

We have picked SQM as the base for our algorithm. The main factors in this decision were that SQM is faster, produces smaller number of triangles, has better edge flow and even without smoothing the generated mesh better resembles the input skeleton. By avoiding the smoothing phase we do not need any input parameters to generate base mesh from an input skeleton. In this chapter we will explain each step of our proposed algorithm as well as extensions like elliptical nodes, cycles, etc.

2.1 Skeleton straightening

Skeleton straightening is a preprocessing step that simplifies bridging of branch node polyhedrons. Straightened skeleton is a skeleton, which nodes in every path between two branch nodes, two leaf nodes, or a branch node and a leaf node are co-linear. Also we have added an extra quality, that angles between child nodes of a branch node should be the same in straightened skeleton, as they are in the input skeleton. To achieve the first condition for each connection node we take the direction of a vector, formed by connection nodes parents position and connection nodes position. The direction vector can be seen in Figure 2.1 as the green arrow. Then we normalize the direction vector and multiply it by the distance between connection node and its child node. The distance is marked by the black curve in Figure 2.1. This vector represents the offset from connection nodes position, at which lies the straightened position of its child node. We then calculate rotation between connection nodes child original position and its new position, in respect to connection nodes position. Finally we rotate all descendants of the connection node. In order to conform to the second condition,

at each branch node we do not alter the position of its child nodes.



Figure 2.1: Skeleton straightening; Left: input skeleton, green arrow represents the direction from connection nodes parents position to connection nodes position, black curve marks the distance between connection node and its child; Right: straightened skeleton

2.1.1 Skinning matrices

In final vertex placement we need to undo the rotations applied to the input skeleton during straightening. We have decided that the best solution is to use skinning, since it can be implemented on GPU and we wanted to move all post-processing on the GPU. Straightened skeleton represents bind pose for skinning purposes and the input skeleton represents reference pose. Now we can calculate skinning matrices required to transform bind pose to reference pose. Traditionally that would require to find the rotation between two corresponding nodes in respect to they parent. Rotating all child nodes in bind skeleton using the same rotation and propagate the rotation calculation to child nodes. But since we know precisely how bind pose was constructed, we can exploit this knowledge and avoid the rotation of child nodes. In fact we do not even need the bind skeleton itself.

This can be seen in Figure 2.1.1. We want to calculate the rotation that would transform circle node to reference pose. We know that circle nodes parent square node is already in reference pose. We also know, that bind pose was constructed in such a way that all connection nodes chldes are co-linear and preserve the distances between nodes. That means from squares reference pose we can calculate, where would be circle node, if we would apply on it the same transformation matrices as were applied to square node. The distance between square and circle node remains constant in both poses. And the direction at which the circle

node would be is the same as the direction from triangle node to square node, which is marked as green arrow in Figure 2.1.1. Now we only need to store the rotation between calculated circle node position, green circle in Figure 2.1.1 and its actual position red circle, with respect to its parent red square node.




Figure 2.2: Rotation estimation from reference pose. Circle: node which rotation we want to estimate; Square: parent of circle node; Triangle: parent of square node; Blue arrows: edges in reference pose; Green circle: circle node position after applying squares skinning matrices; Green arrow: direction from square to green circle node;

2.2 BNP generation

For each branch node we calculate the direction from said branch node to each of its children and to its parent. Then we create rays, with origin in branch nodes position and direction corresponding to the direction calculated previously. These rays can be seen in Figure 2.2 (a) as blue arrows. We calculate the intersection of each of these rays with a sphere associated with the branch node. We store each intersection in a set of intersection points. Now we triangulate the intersection points. Different algorithms can be used to achieve the same effect, but we have picked Delaunay triangulation in spherical domain, which was also used in the original paper. The algorithm works like standard Delaunay triangulation, but the predicate deciding whatever newly inserted triangle would lie in the circle of an already existing triangle is replaced. The new predicate compares angle between the newly inserted triangles normal with normals of already existing triangles. Result of triangulation is shown in Figure 2.2 (b) as the blue triangle. The generated polyhedron is now subdivided by inserting a point

in the center of each face and in the middle of each edge. The vertex inserted in the center of each face is then connected with all vertices corresponding to the same face. So each triangle is subdivided into six smaller triangles. The newly inserted points are then projected onto the sphere associated with the node. The subdivision and projection is necessary, because otherwise polyhedrons that would be generated with co-planar, or nearly co-planar intersection points would have no volume or very little volume respectively. To project the newly inserted vertices onto the sphere, we once again use a ray-sphere intersection. The origin of the ray is the position of each newly inserted vertex. The direction of the ray is mean normal of the faces that are connected with the vertex. This means that for the vertices in the center of each face the normal of the subdivided face is used. For vertices inserted in the middle of each edge the mean normal of faces corresponding to that edge is used. Final polyhedron is shown in Figure 2.2 (c).




Figure 2.3: BNP generation process. (a) green is a branch node, blue arrows represent direction vectors of rays, red circles represent child nodes; (b) blue triangle is the result of triangulation; (c) final subdivided BNP;

Triangulation of intersection vertices can sometimes create obtuse triangles. These are problematic, because when we insert the vertex in the middle of an obtuse triangle the one-rings of intersection vertices are not convex. That is if we would project them onto a plane defined by their principal axis and one of the vertices, the resulting polygon would not be convex. In Figure 2.2 we can see on the left how a polyhedron looks when it is generated with obtuse triangles. Expected central vertex position is marked by red arrow, also expected edges are marked as yellow lines. This is not desirable, since it would cause problems during

cyclic mesh generation. To remedy this situation we calculate the projection of the central vertex in a different manner. The origin of its ray is still vertex position. But we use a new direction vector. This vector is the normal of a triangle formed by the projected points that were inserted in the middle of each edge. The result can be seen in Figure 2.2 on the right.




Figure 2.4: Obtuse triangle problem. Left: polyhedron with obtuse triangle. Red arrow marks vertex expected position and yellow lines mark expected edges; Right: polyhedron after applying our fix;

2.3 BNP refinement

During the penultimate step of the algorithm, BNP joining, we want to connect two BNPs, with tube consisting solely from quadrilaterals. To ensure that we can use quadrilaterals only, we want intersection vertices connected via a path, to have the same number of nodes in their one-ring neighbourhood. We take the notion of Link Intersection Edge (LIE) from [9]. A LIE is a set of edges that are between vertices of two intersection vertices one-ring neighbourhoods, as can be seen in Figure 2.3 (a). One LIE is represented with yellow colored edges and second LIE is represented with red colored edges. The refining of one BNP is a one-pass procedure.

Preprocessing We start the refining of a BNP by creating a map of LIEs corresponding to each intersection vertex. For each intersection vertex we store its corresponding LIEs, as well as the number of splits required by that vertex, to have the same valency as its corresponding intersection vertex. An intersection vertex can be connected with three types of vertices.

Each type defines how much we can split LIEs corresponding to that intersection vertex.

- **Parents intersection vertex** - the number of splits is fixed. That is after splitting the valency of these two vertices must be exactly the same. This is necessary otherwise splits in child BNP could cause splits in parents BNP, which could result in infinite loops.
- **Branch node intersection vertex** - difference between valencies of the two vertices is calculated. If the number is negative, that is corresponding vertex has smaller valency, we prefer not to split LIEs corresponding the vertex any more. If the number is positive, that is corresponding vertex has higher valency, we split LIEs corresponding to the vertex so that the vertex has at least the same valency as its corresponding vertex.
- **Leaf node** - corresponding LIEs can be split as much has needed.




Figure 2.5: BNP refinement. Yellow edges represent Link Intersection Edge (LIE) corresponding to second intersection edge of the BNP. Red edges represent LIE corresponding to third intersection edge. Yellow circle marks the first intersection vertex. BNP before subdivision (a); BNP after one subdivision and one smoothing step (b); Final BNP after two subdivisions and two smoothing steps (c);

Subdivision We loop through all intersection vertices, starting with the intersection vertex connected with parents BNP and continuing with the rest. In each cycle, we split LIEs corresponding to current intersection vertex, until the valency of the intersection vertex is equal to its required number of splits. Every time we want to split a LIE we heuristically select the best, based on two filters. The first filter is the splitting required by the other

intersection vertex belonging to the same LIE. If two or more LIEs require the same number of splits, we use a second filter. This filter picks the LIE that was split the least. When we are splitting a LIE, we always split a representative edge, which is the first edge of the LIE. Because of that we need to apply a smoothing scheme to roughly equalize the lengths of edges in a LIE.

The whole process can be seen in Figure 2.3. In Figure 2.3 (a) the yellow circle marks an intersection vertex that needs to be split twice. Yellow edges represent a LIE corresponding to second intersection vertex and red edges represent a LIE corresponding to third intersection vertex. Both LIEs have the same need to be split and none of them was split previously. For the first split the yellow LIE is selected, subdivided once and smoothed. The result of the first split is shown in Figure 2.3 (b). For the second split the need of both LIEs is still the same. But yellow LIE was already split, so this time the red LIE is selected, subdivided and smoothed. Final refined BNP is shown in Figure 2.3 (c).

2.3.1 Smoothing

Since we always split only one representative edge of a LIE, we are applying a smoothing scheme, to equalize the length of edges in a subdivided LIE. The smoothing is very important because directly on it depends the quality of the generated mesh. Ideally the length of each edge in a smoothed LIE would be equal. However since smoothing is applied after every subdivision, the smoothing algorithm should to be reasonably fast. We propose four smoothing schemes. These smoothing schemes are illustrated in Figure 2.3.1, where the polyhedron from Figure 2.3 is subdivided twice and then smoothed with various smoothing schemes.

Averaging smoothing New position for each vertex in a LIE is calculated by averaging one-ring neighbourhood corresponding to the vertex. We start with the last vertex of a LIE, that is the vertex on the last edge of a LIE and move towards the first vertex. We move each vertex, except the first and the last vertices, to the barycentre of its one-ring neighbourhood and project them back onto the sphere corresponding to BNPs node. The resulting smoothed polyhedron is shown in Figure 2.3.1 (a). This approach is iterative and would need several iteration to achieve global optimum, however we have found that one iteration is enough for our needs.

Quaternion smoothing For each LIE we calculate a quaternion representing the rotation from the first vertex of the LIE to the last vertex of the LIE. From each quaternion we extract its corresponding axis of rotation and angle of rotation. We smooth only points between the first and the last vertex, so the calculated axis of rotation and angle are constant. During each smoothing step, we first count the number of vertices in a LIE. Then we divide the angle of rotation by that number and form a new quaternion from already calculated axis of rotation and the newly calculated angle. For each vertex in a LIE between first and last we apply the rotation stored in the quaternion and update its position. This method produces LIEs, that lie on small circles of their corresponding sphere. The spacing between vertices is regular and thus its very suitable for our needs. The result of quaternion smoothing is shown in Figure 2.3.1 (b).

Area weighted Laplacian For Laplacian smoothing we have adapted the algorithm described in [7]. The weights used for smoothing are based on the one-ring area of each vertex. We use one iteration of Laplacian smoothing and then project the new vertices back onto their corresponding sphere. The result is shown in Figure 2.3.1 (c). The result is usable, since the edges have better distributed length, than without smoothing. But averaging and quaternion smoothing produce results of better quality.




Figure 2.6: LIE smoothing schemes. Shows results after applying averaging smoothing (a); Quaternion smoothing (b); Area weighted Laplacian smoothing (c); and Valency weighted Laplacian smoothing (d);

Valency weighted Laplacian We use the same algorithm as for area weighted Laplacian, but we use different weights for vertices. The weights depend on the valency of each vertex.

After Laplacian smoothing vertices are projected back onto their corresponding sphere. The result of this smoothing is very similar to area weighted Laplacian and is shown in Figure 2.3.1 (d).

2.4 BNP joining

After refinement of BNPs, each intersection vertex has the same valency as its corresponding intersection vertex. Now BNPs can be joined by tubes, consisting from quadrilaterals only. We loop through each branch node in a depth-first search from skeletons root. We process each BNP in the following manner. We start with the whole BNP Figure 2.4 (a). We loop through all intersection vertices corresponding to current BNP. We remove each intersection vertex and their corresponding faces and edges from current BNP. In Figure 2.4 (b) we can see the removal of third intersection vertex, after first and second intersection vertices were joined. After the removal of an intersection vertex we continue joining all nodes on the path that produced the removed intersection vertex. We take the vertices forming the former one-ring of the removed intersection vertex. For each connection node on the path we duplicate the one-ring vertices, translate them to the nodes position and project them onto the sphere associated with the connection node. For the translation we construct a plane. The origin of the plane is the position of the connection node. The normal is the vector from the branch node to the connection node. Each vertex is translated along the normal until it lies on the plane. The projection is done by normalizing the vector from connection nodes position to each vertex, and multiplying it by the radius of connection nodes corresponding sphere. These new vertices are then connected with previous set of vertices with faces and they are passed as current one-ring vertices for joining. In Figure 2.4 (c), we can see the translated connection node vertices, before projection was applied. In Figure 2.4 (d), the duplicated one-ring vertices are projected onto their associated sphere. After going through all connection nodes we can end either in a branch node or in a leaf node.

Branch node termination We start by removing the destination branch node corresponding intersection vertex and faces and edges connected to it. The tube generated from connection nodes should be joined with destination intersection vertex former one-ring. For this we need to find two corresponding vertices, which would form the first pair and define the pairing of the rest of the vertices. We can use multiple methods to find the first pair. The choice




Figure 2.7: BNP joining process. Polyhedron before joining (a); Polyhedron with removed faces corresponding to an intersection vertex (b); New vertices for connection node before projection (c); Projected vertices of connection node (d);

of the method affect output base mesh quality, as a wrong selection can lead into twists in the base mesh. We tested three methods.

- **Minimal Euclidean distance** - An arbitrary vertex is selected from current one-ring and its euclidean distance to each vertex of destination one-ring is calculated. We match the vertex with the closest one. The main disadvantage of this method is that it is greedy and ends in local minima. That is to have the best overall mesh quality a more distant vertex would be a better choice. This can be seen in Figure 2.4 (a), where some connecting faces are twisted.
- **Minimal angle** - An arbitrary vertex is selected from the current one-ring. The direction from the selected vertex to each vertex of destination one-ring is calculated. Then we calculate the angle between this direction and the direction from original branch node to destination branch node. As a match we pick a vertex, where the calculated angle was minimal. As before this method is greedy and can fall into local minima. An example of this is shown in Figure 2.4 (b), where some faces connecting the hand are twisted.
- **Minimal total euclidean distance** - An arbitrary vertex is selected from current one-ring. We pair the vertex with each vertex of destination one-ring. Each pair defines a mapping from the current one-ring vertices to destination one-ring vertices. We calculate the euclidean distance between all paired vertices and store it as total euclidean distance. Then we pick as pair the vertex where the total distance was minimal. The disadvantage of this approach is that it takes longer to compute, but it avoids local

minima, by testing all the possible solutions and picking the best one. Since one-rings usually consist of few vertices, we have decided to use this method. The resulting mesh using this pairing is shown in Figure 2.4 (c). Among the possible results in Figure 2.4 (c) has minimal twisting compared to (a) and (b).

When the pairing of vertices is defined, new edges can be formed between each pair and connect the geometry.




Figure 2.8: Minimal euclidean distance (a); Minimal angle (b); Minimal total Euclidean distance (c);

Leaf node termination As in original paper [9], we have decided to terminate leaf nodes with triangular faces. Each face has two vertices from the current one-ring and one vertex is the position of the leaf node.

2.5 Final vertex placement

We use skinning matrices calculated during skeleton straightening to reverse the rotations applied during straightening. Various methods can be used to rotate vertices back into their original position, but we have found that Linear Blend Skinning as described in [5] is enough for our needs. Skinning transformation can be applied directly on GPU in vertex or tessellation shaders. We have also implemented Linear Blend Skinning on CPU. This is advantageous for comparison purposes, as well as when we are generating base meshes from cyclic skeletons. Since we are using triangulations on CPU to close generate cyclic skeletons, we need to apply Linear Blend Skinning before the triangulations. Skinning could be still calculated on GPU and the resulting geometry could be send back to CPU for processing, however

we have found that our generated base meshes are relatively simple and the overhead from transmitting data between CPU and GPU is unnecessary.

2.6 Ellipsoid Nodes

An ellipsoid can be defined as a sphere with associated transformation matrix. We take advantage of this representation of ellipsoids. Instead of more complex ray-ellipsoid intersection that would have to be computed at each ellipsoid node, we have decided to split each ellipsoid node into a sphere and a transformation matrix. First our base mesh algorithm is evaluated as described in Chapter 2 with spherical nodes. After that we send the transformation matrices corresponding to each ellipsoid node to GPU. The vertices corresponding to each ellipsoid node are transformed directly in vertex shader. Thanks to this ellipsoid nodes require minimal extra computing resources from CPU. The results can be seen in Figure 2.9. An input skeleton with ellipsoid nodes is displayed in Figure 2.9 (a). The generated base mesh is shown in Figure 2.9 (b) and from different camera angle in Figure 2.9 (c).




Figure 2.9: Ellipsoid nodes. Skeleton with ellipsoid nodes specified (a); Base mesh generated from skeleton (b); Base Mesh from different angle (c);

2.7 Tessellation

Tessellation shaders available since OpenGL 4.0 are used to tessellate the generated base mesh. Two connected spherical nodes, a parent and a child, implicitly define a truncated cone between them. The base of the cone has the radius of parent spherical node and the top of the truncated cone has the radius of child spherical node. Each vertex generated during tessellation is projected onto this cone. The projection is done by translating the vertex along

its normal until it reaches the surfaces of the cone. Tessellation is shown in Figure 2.10, the generated base mesh is shown in (a) and the tessellated base mesh in (b). However during this step the generated base mesh gains volume and the newly generated vertices can intersect the tessellated base mesh. This effect can be seen in Figure 2.10 (c). To recover from this situation, we detect sharp vertices in the input mesh and apply a smoothing scheme. Sharp vertices are vertices which faces are forming sharp angles. In tessellation shader we have access only to one patch and its vertices. So we compute the sharpness of each vertex by comparing and thresholding the normal of each vertex with the normal of the patch. We apply bezier smoothing to modify the radius of the truncated cone. The smoothed mesh is shown in Figure 2.10 (d). Currently the smoothing bezier curve is constant, but it could be dynamically changed based on the sharpness of the vertices.




Figure 2.10: Tessellation. Not tessellated mesh (a); Tessellated mesh with 20 subdivisions (b); Tessellated mesh with self intersection (c); Tessellated mesh with smoothing (d);

2.8 Capsule Ending

Generating of capsules can be approached in two ways. The first is to generate a capsule at each leaf node corresponding to its radius. The second is inserting additional nodes into the

input skeleton, with decreasing radius, that would approximate a capsule. We have implemented the second approach, because it fits nicely into our pipeline. Capsules generated this way, can be directly tessellated on the GPU, without any additional processing. At each capsule leaf node, we insert additional nodes into the input skeleton, proportional to the radius of the capsule node. This transformation is shown in Figure 2.11, where the input skeleton (a) transforms into (b). The radius of each node is decreased according to Equation 2.1, where $nodeRadius$ is the radius of capsule node and $step$ is a number between $(0 - 1]$, that represents the distance from center of the capsule to its edge. Final tessellated capsule is shown in Figure 2.11 (c).

$$newRadius = \sqrt{nodeRadius^2 * (1 - step^2)} \quad (2.1)$$




Figure 2.11: Capsule generation. Input skeleton (a); Skeleton with generated extra nodes (b); Tessellated base mesh with capsule (c);

2.9 Linear Skeletons

Linear skeletons without branch nodes lack the initial geometry that is generated during BNP generation step described in Section 2.2. Additional nodes could be inserted into the input skeleton, to form at least one branch node, but we have found that it needlessly disturbs the flow of the output mesh. Instead we decided to use a different approach. We introduce an additional input parameter N , which specifies how many vertices should be generated for each node of the linear skeleton. This parameter does not decrease the robustness of our approach, because additional vertices are generated during tessellation and the original number of vertices is negligible.

First step of the algorithm is setting the root to be the head of the input linear skeleton. Next step of the algorithm is straightening of the input linear skeleton. The input skeleton is shown in Figure 2.12 (a). Next, N vertices are generated around first connection node, which is a child of the root node. These vertices are distributed regularly around the node, by slerping a quaternion, which center of rotation is nodes position, axis of rotation is the direction from connection node to root node and magnitude is $360/N$. Newly generated vertices are then joined with other vertices as in original base mesh algorithm. Leaf nodes form a triangle fan and connection nodes form a tube of quadrilaterals. The joined linear base mesh is shown in Figure 2.12 (b). Skinning matrices are used to transform the generated linear skeleton into its input pose Figure 2.12 (c).




Figure 2.12: Linear base mesh generation. Input linear skeleton (a); Straightened and joined linear skeleton (b); Final linear base mesh (c);

2.10 Root That Is Not a Branch Node

If the root of the input skeleton is not a branch node and a branch node is present in the skeleton, we can find it with depth first search. When we have at least one branch node we can re-root the tree, so that the located branch node would be the root of the tree. We apply Algorithm 1 to re-root the skeleton. We want to make the input node the root of the skeleton. For that it has to have no parent node and all other nodes in the skeleton should have a parent node. We set the new root to have no parent and store its former parent as *current* node and the parent of *current* node as *currentParent* node. In a loop we set *currentParent* as a child of *current* node. In order to advance along the skeleton, we set *current* node to the parent of *current* node and *currentParent* node to the parent of *currentParent* node. This way we

ascend along the skeleton, until we reach the former root node and the algorithm terminates.

Algorithm 1 ReRoot

Input: *node*{the node that should be the new root}

current \leftarrow *node.parent*

currentParent \leftarrow *current.parent*

node.parent \leftarrow NULL

current.parent \leftarrow *node*

current.removeChild(node)

node.addChild(current)

while *current* \neq *root* **do**

current.addChild(currentParent)

parent.removeChild(current)

newParent \leftarrow *currentParent.parent*

currentParent.parent \leftarrow *current*

current \leftarrow *currentParent*

currentParent \leftarrow *newParent*

end while

root = *node*

2.11 Cyclic Skeletons

Our last improvement is generation of base meshes from cyclic skeletons. The cycle can be placed anywhere in the input skeleton. The base algorithm could not be modified to allow generation of cyclic meshes, because during BNP refinement step of the algorithm a cycle could cause an infinite loop. However we can modify the input skeleton in a way, that would allow us to generate cyclic skeletons. As the input we have a cyclic skeleton Figure 2.14 (a). Cyclic edge is marked with yellow color and cyclic nodes with green and violet colors.

First, we split the cycle, by removing the yellow cyclic edge. To each cyclic node we add an extra child node as shown in Figure 2.14 (b). Light green node for green cyclic node and pink node for violet cyclic node. These new nodes serve to preserve the skinning matrices, that will rotate tubes generated from cyclic nodes, to face each other. This can be seen in

Figure 2.14 (c). Base mesh was generated as described in Chapter 2, with one exception. The triangles that should be generated for light green and pink nodes were omitted. Now the gap between cyclic nodes can be closed. We first project vertices associated to each cyclic node to a plane with origin at $O(0, 0, 0)$ and normal $n(0, 1, 0)$. Next, we normalize the vertices so that vertices associated with violet node lie at a circumference with radius 1 and vertices associated with green node lie at circumference with radius 2. The position of projected points is shown in Figure 2.13 (a), where vertices associated with green node have green color and vertices associated with violet node have violet color. Now we execute a Delaunay triangulation on the transformed points. After the triangulation is done, we exclude triangles generated solely between green or violet vertices. The remaining triangles represent the faces that should be generated in order to close the gap between cyclic nodes, as can be seen in Figure 2.13 (b). Final cyclic mesh is shown in Figure 2.14 (d).




Figure 2.13: Delaunay triangulation of input points (a) used to generate faces between vertices forming the cycle (b); Green points represent vertices corresponding to cycle node 1 and violet represent vertices of cycle node 2 from Figure 2.14




Figure 2.14: Cyclic skeleton base mesh generation. Cyclic skeleton, cyclic edge marked with yellow color, cyclic nodes with green and violet (a); Split cycle with one inserted node for each former cyclic nodes lightgreen for green node and pink for violet node (b); Generated base mesh before the cycle is closed (c); Generated base mesh after the cycle is closed (d);

Chapter 3

Implementation

We have implemented a system, where users can load and edit skeletons from files, or start by scratch from one skeletal node. Our system is based on model-view-controller pattern. Model is stored in *BMMAlgorithm* and *BMMNode* classes. These classes take care of storing the skeleton, exporting skeleton to different formats and executing the base manifold mesh algorithm. Controller is represented by *BMMController* class. This class prepares data from model for visualization and handles input from view. View consists of two parts. First, is *OpenGL* window, which displays data provided from controller. Second, is a Windows form, that provides Graphical User Interface (GUI) elements, which serve to make input easier.

In this chapter we will describe how we implemented our solution. We will start with used programming language, Integrated Developer Environment (IDE), tools and libraries. Next, we will describe each part of model-view-controller pattern in greater detail. After that we will show some programmable shaders used in our implementation.

3.1 Programming Language, IDE, Tools, and Libraries

We have selected C++ as programming language. The main reasons are that C++ is fast and well established programming language. That means that most of the open source libraries are available for C++ and many of them are available exclusively for C++. As IDE we are using Visual Studio 2012, which was the newest version of Visual Studio, when we started our work. Visual Studio is a powerful and well established IDE for developing on Windows operating systems. We have used nVidia nSight for Visual Studio 2012 that is a debugger for

graphics cards, which allows to set breakpoints in shaders during execution. We have also used several open source libraries in our project. We will briefly describe the key libraries. To program the GPU we use OpenGL Shading Language (GLSL) version 4.3.

- **Boost [1]** Boost is a set of libraries encapsulating various task like basic input output system, smart pointers, serialization, matrices, etc. We use Boost primary for class serialization. Thanks to boost we can serialize classes into XML files, which can be stored and loaded from disk. This XML files can even be shared between different applications, so we can import skeletons from third party programs.
- **OpenMesh [10]** Is an open source half edge data structure library. We use it to store all meshes in our applications. OpenMesh library is capable of storing meshes composed solely from triangles as well as meshes composed of arbitrary polygons. OpenMesh has pre-calculated many convenient iterators for one-ring neighbourhoods, faces and edges adjacent to a vertex, etc. The library is also capable of exporting stored meshes into Wavefront .obj files.
- **GLEW [3]** The OpenGL Extension Wrangler Library (GLEW) provides efficient run-time mechanisms for activating OpenGL extensions. We use this library to expose OpenGL 4.3 functionality to our application.
- **OpenGL [4]** Open source cross platform application programming interface. We use it to leverage the computation capability of GPUs to achieve hardware accelerated rendering.
- **GLM [8]** OpenGL Mathematics library that provides the same vector and matrix operations as available in GLSL. GLM also provides several function, that are deprecated in newer versions of OpenGL. We use this library to manipulate camera, model matrices, and for general matrix computations.

3.2 Model

3.3 View

3.4 Controller

3.5 GPU Shaders

Appendix A

T_EX

L^AT_EX, T_EX

Bibliography

- [1] Boost. Boost C++ libraries. <http://www.boost.org/>, 2014. [Online; accessed 24-February-2014].
- [2] Michael Mc Donnell. Skeleton-based and interactive 3d modeling. Master’s thesis, Technical University of Denmark, 2012.
- [3] GLEW. The OpenGL extension wrangler library. <http://glew.sourceforge.net/>, 2014. [Online; accessed 24-February-2014].
- [4] OpenGL Khronos Group. The industry’s foundation for high performance graphics. <http://www.opengl.org/>, 2014. [Online; accessed 24-February-2014].
- [5] Ladislav Kavan, Steven Collins, Jiri Zara, and Carol O’Sullivan. Skinning with dual quaternions. In *2007 ACM SIGGRAPH Symposium on Interactive 3D Graphics and Games*, pages 39–46. ACM Press, April/May 2007.
- [6] Zhongping Ji Ligang Liu and Yigang Wang. B-mesh: A fast modeling system for base meshes of 3d articulated shapes. *ComputGraphForum*, 29(7):2169–77, 2010.
- [7] Martin Madaras. Todo. Master’s thesis, Comenius university of Bratislava, TODO.
- [8] OpenGL Mathematics. A C++ mathematics library for graphics programming. <http://glm.g-truc.net/0.9.5/index.html>, 2014. [Online; accessed 24-February-2014].
- [9] J. A. Bærentzen M. K. Misztal and K. Welnicka. Converting skeletal structures to quad dominant meshes. *Computers & Graphics*, 36(5):555–561, 2012.
- [10] OpenMesh. OpenMesh open half edge data structure. <http://openmesh.org/>, 2014. [Online; accessed 24-February-2014].